# BEFORE WE GET STARTED

Have a Github account:
https://github.com/join

Download the Github desktop client
https://desktop.github.com

Clone the today's materials from Github (green button "Clone or download")

https://github.com/Dt1431/chi-ds-5-lesson-3

My email: daveturner81@gmail.com
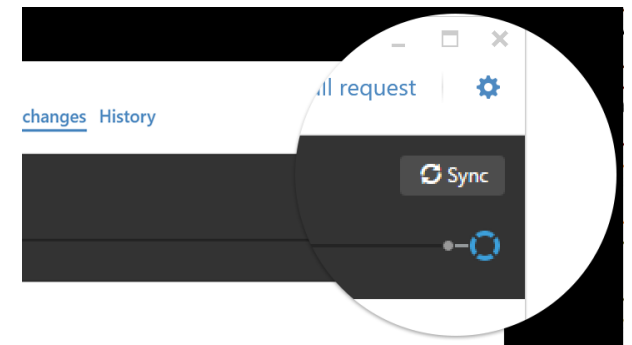
# STATISTICS FUNDAMENTALS

*David Turner*

*CEO, Waitbot Inc.*

# LEARNING OBJECTIVES

‣ Use NumPy and Pandas libraries to analyze datasets using basic summary statistics: mean, median, mode, max, min, quartile, inter-quartile range, variance, standard deviation, and correlation

‣ Create data visualizations - including: line graphs, box plots, and histograms- to discern characteristics and trends in a dataset

‣ Identify a normal distribution within a dataset using summary statistics and visualization

‣ ID variable types and complete dummy coding by hand

# DUE: 3/27 @ Midnight

Project: Using a dataset of your choice, create a Jupyter Notebook or python script that demonstrates your understanding of python, pandas, and statistics (this weeks topic).

- ‣ Requirements:
  - ‣ Import and clean data
  - ‣ Create calculated columns (e.g. dummy variables, new groupings)
  - ‣ Describe your data (statistics, visualizations)
  - ‣ Identify areas for further research
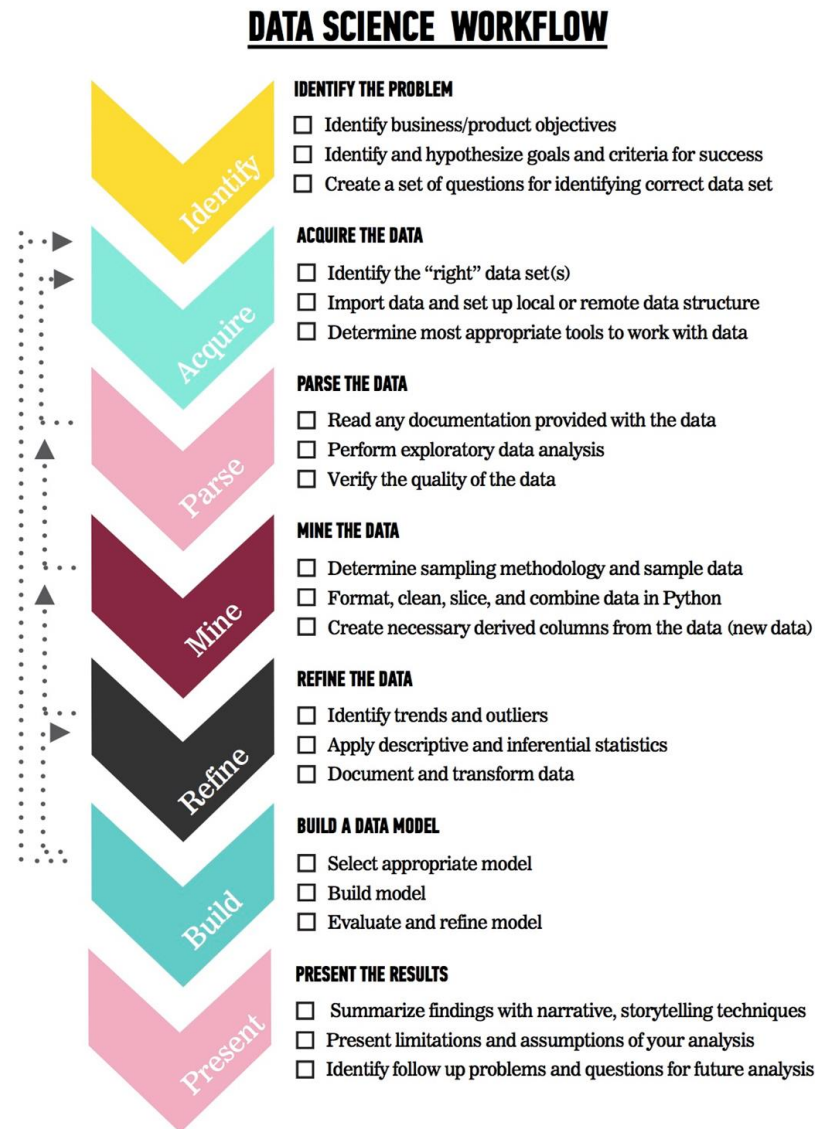  - ‣ Bonus: Combine with a second dataset

# STATISTICS FUNDAMENTALS

# LET'S REVIEW THE DATA SCIENCE WORKFLOW

The steps:

1. Identify the problem
2. Acquire the data
3. Parse the data
4. Mine the data
5. Refine the data
6. Build a data model
7. Present the results

## DATA SCIENCE WORKFLOW

**IDENTIFY THE PROBLEM**
- ☐ Identify business/product objectives
- ☐ Identify and hypothesize goals and criteria for success
- ☐ Create a set of questions for identifying correct data set

**ACQUIRE THE DATA**
- ☐ Identify the "right" data set(s)
- ☐ Import data and set up local or remote data structure
- ☐ Determine most appropriate tools to work with data

**PARSE THE DATA**
- ☐ Read any documentation provided with the data
- ☐ Perform exploratory data analysis
- ☐ Verify the quality of the data

**MINE THE DATA**
- ☐ Determine sampling methodology and sample data
- ☐ Format, clean, slice, and combine data in Python
- ☐ Create necessary derived columns from the data (new data)

**REFINE THE DATA**
- ☐ Identify trends and outliers
- ☐ Apply descriptive and inferential statistics
- ☐ Document and transform data

**BUILD A DATA MODEL**
- ☐ Select appropriate model
- ☐ Build model
- ☐ Evaluate and refine model

**PRESENT THE RESULTS**
- ☐ Summarize findings with narrative, storytelling techniques
- ☐ Present limitations and assumptions of your analysis
- ☐ Identify follow up problems and questions for future analysis

Identify
Acquire
Parse
Mine
Refine
Build
Present

# TODAY

‣ We're going to begin to talk about step 3:  Parsing the Data

‣ We'll begin to talk about the fundamentals of Statistics

# LAYING THE GROUND WORK

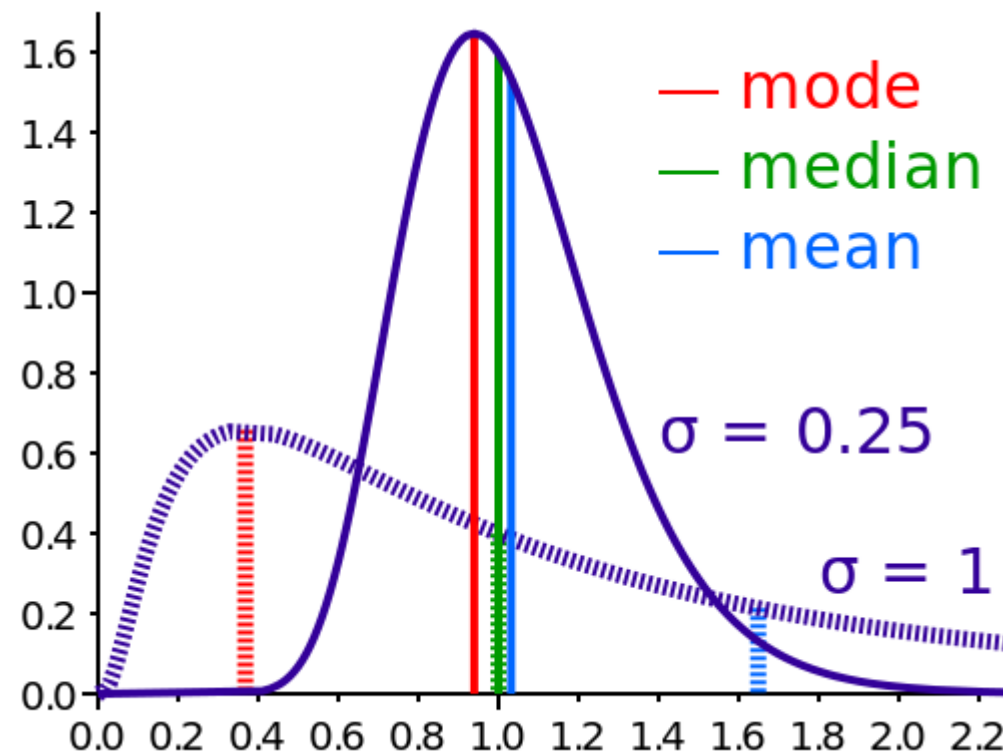# WE'RE GOING TO COVER SEVERAL TOPICS

‣ Mean
‣ Median
‣ Mode
‣ Max
‣ Min
‣ Quartile
‣ Interquartile Range
‣ Variance
‣ Standard Deviation
‣ Correlation

# MEAN

‣ The mean of a set of values is the sum of the values divided by the number of values. It is also called the average.

$$\overline{X} = \frac{\sum X}{N}$$

# MEAN EXAMPLE

‣ Find the mean of 19, 13, 15, 25, and 18.

# MEAN EXAMPLE

▸ Find the mean of 19, 13, 15, 25, and 18.

$$\frac{19 + 13 + 15 + 25 + 18}{5} = \frac{90}{5} = 18$$

# MEDIAN

‣ The median refers to the midpoint in a series of numbers.

‣ To find the median

   ‣Arrange the numbers in order smallest to largest.

   ‣If there is an odd number of values, the middle value is the median.

   ‣If there is an even number of values, the average of the middle two values is the median.
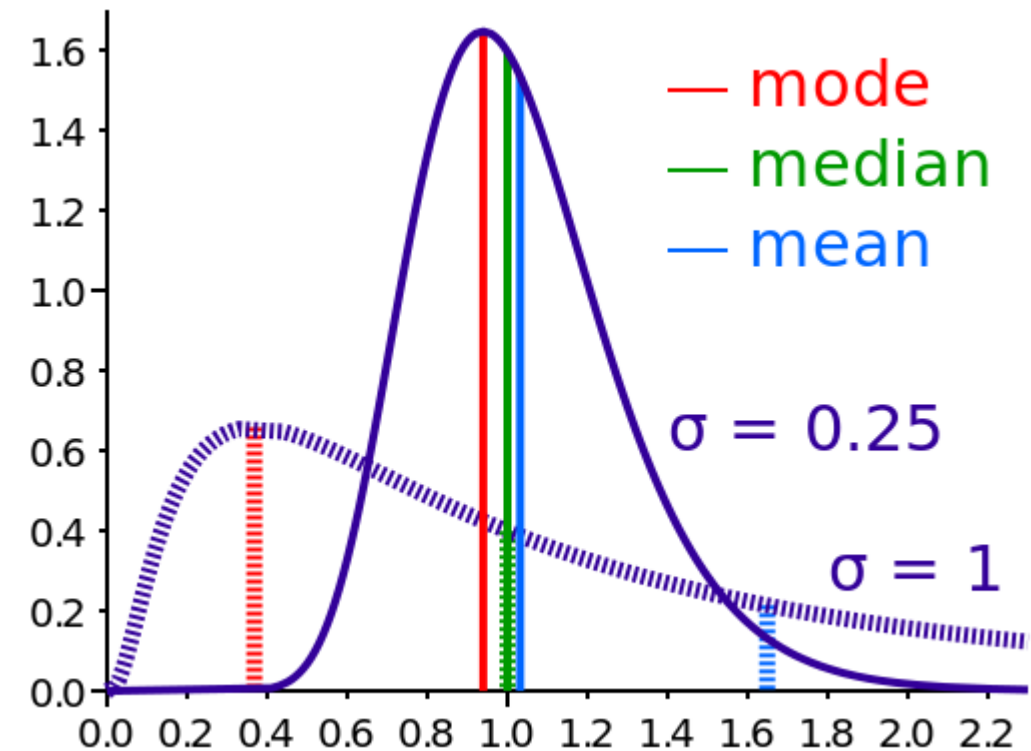
# MEDIAN EXAMPLE

‣ Find the median of 19, 29, 36, 15, and 20.

# MEDIAN EXAMPLE

‣ Find the median of 19, 29, 36, 15, and 20.

Ordered Values:

15, 19, 20, 29, 36

20 is the median

# MEDIAN EXAMPLE

▸ Find the median of 67, 28, 92, 37, 81, 75.

# MEDIAN EXAMPLE

‣ Find the median of 67, 28, 92, 37, 81, 75.

Ordered Values:

28, 37, 67, 75, 81, 92

67 and 75 are the middle values.

$$\frac{67 + 75}{2} = \frac{142}{2} = 71$$

71 is the median.

# MODE

‣ The mode of a set of values is the value that occurs most often.

‣ A set of values may have more than one mode or no mode.
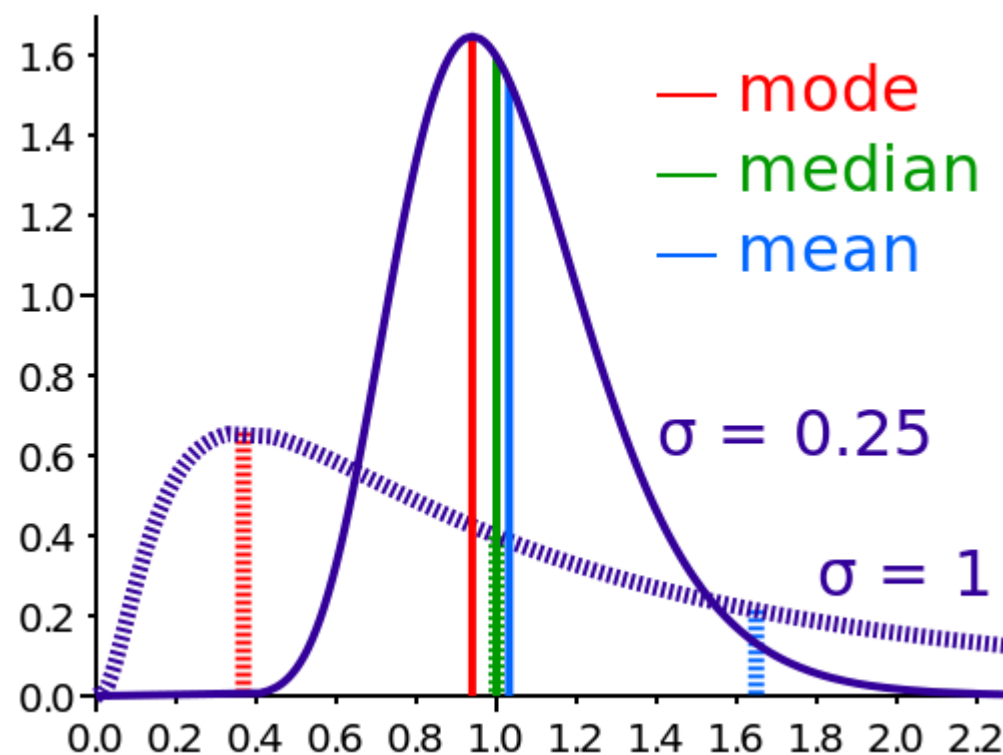
# MODE EXAMPLE

‣ Find the mode of 15, 21, 26, 25, 21, 23, 28, and 21.
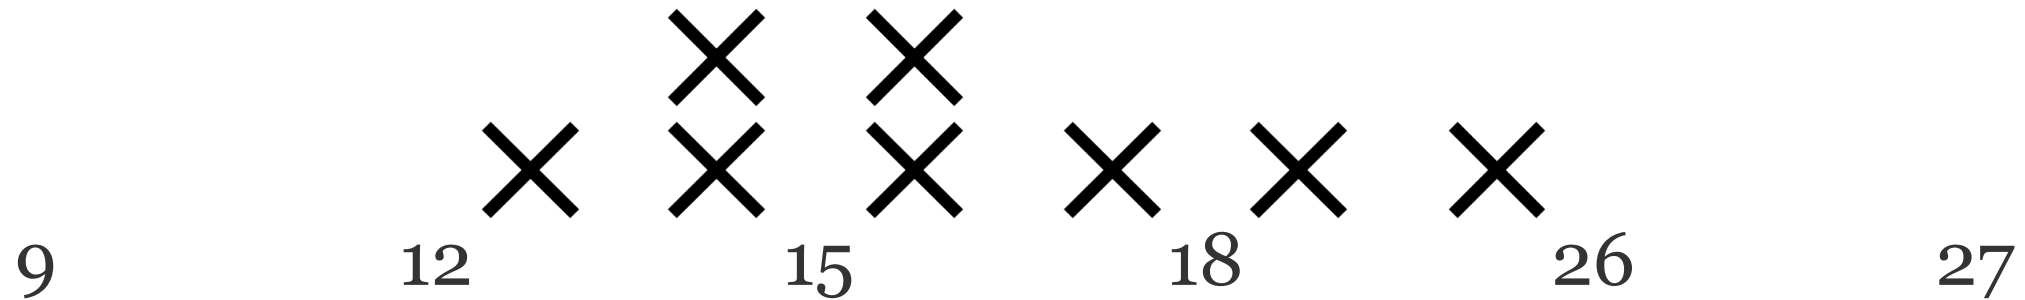
# MODE EXAMPLE

‣ Find the mode of 15, 21, 26, 25, 21, 23, 28, and 21.

X
X
X    X    X    X    X    X

15      21      23      25      26      28
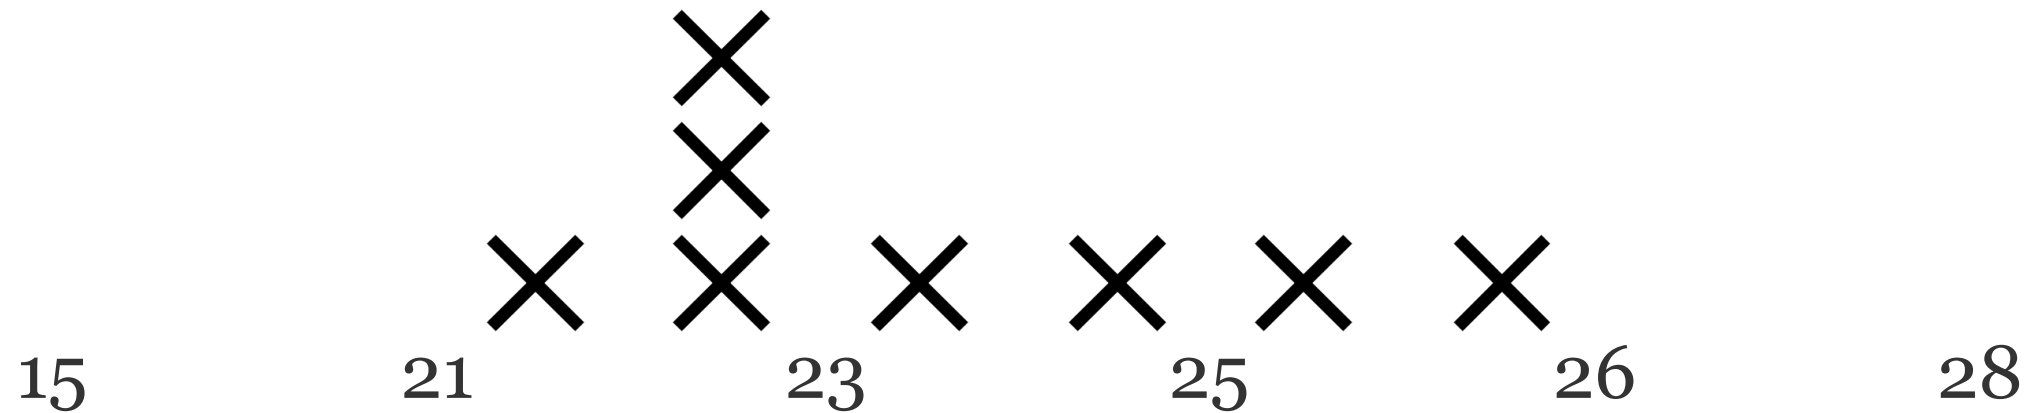
21 is the mode because it occurs most frequently

# MODE EXAMPLE

‣ Find the mode of 12, 15, 18, 26, 15, 9, 12, and 27.

# MODE EXAMPLE

‣ Find the mode of 12, 15, 18, 26, 15, 9, 12, and 27.



12 and 15 are the modes since the both occur twice.

# MODE EXAMPLE

‣ Find the mode of 4, 8, 15, 21, and 23.

# MODE EXAMPLE

‣ Find the mode of 4, 8, 15, 21, and 23.

$$\times \quad \times \quad \times \quad \times \quad \times$$

4        8        15        21        23

There is no mode since all values occur the same number of times.

# ACTIVITY: KNOWLEDGE CHECK

EXERCISE

**ANSWER THE FOLLOWING QUESTIONS (5 minutes)**

1. For the following groups of numbers, calculate the mean, median and mode by hand. Also determine the min and max.

    a. 18, 24, 17, 21, 24, 16, 29, 18

    b. 75, 87, 49, 68, 75, 84, 98, 92

    c. 55, 47, 38, 66, 56, 64, 44, 39

**DELIVERABLE**

Answers to the above questions

# SUMMARY STATISTICS IN PANDAS

# CODEALONG: SUMMARY STATISTICS IN PANDAS

‣ Open the starter-code notebook located in lessons/lesson-03/code/starter-code of the class repo.

# CODEALONG PART 1: BASIC STATS

‣ We can use Pandas to calculate the mean, median, mode, min, and max.

```
Methods available include:

.min() - Compute minimum value

.max() - Compute maximum value

.mean() - Compute mean value

.median() - Compute median value

.mode() - Compute mode value

.count() - Count the number of observations
```
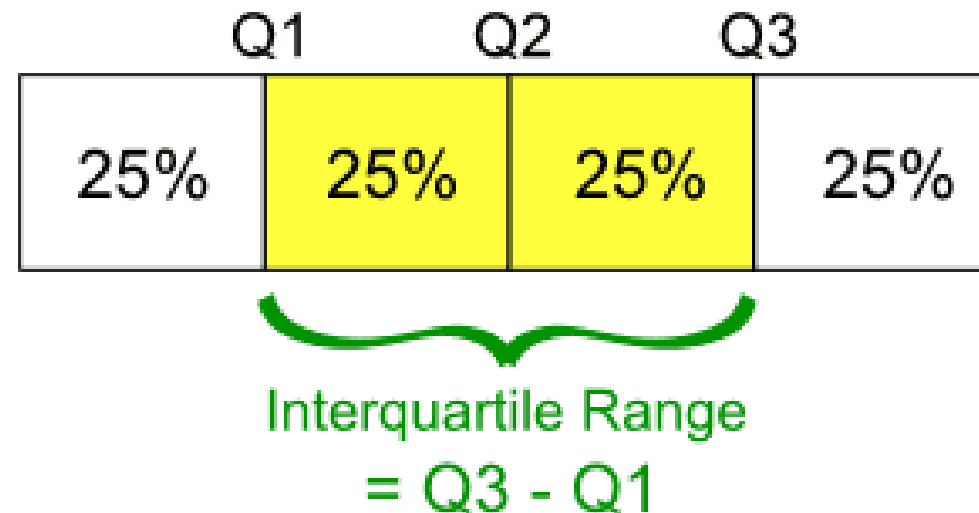
# QUARTILES AND INTERQUARTILE RANGE

‣ Quartiles divide a rank-ordered data set into four equal parts.

‣ The values that divide each part are called first, second, and third quartiles, denoted Q1, Q2, and Q3, respectively.

‣ The interquartile range (IQR) is Q3 - Q1, a measure of variability.

|  | Q1 | Q2 | Q3 |  |
|---|---|---|---|---|
| 25% | 25% | 25% | 25% | |

Interquartile Range
= Q3 - Q1

# CODEALONG PART 2: BOX PLOT

‣ Box plots give a nice visual of min, max, mean, median, and the quartile and interquartile range.
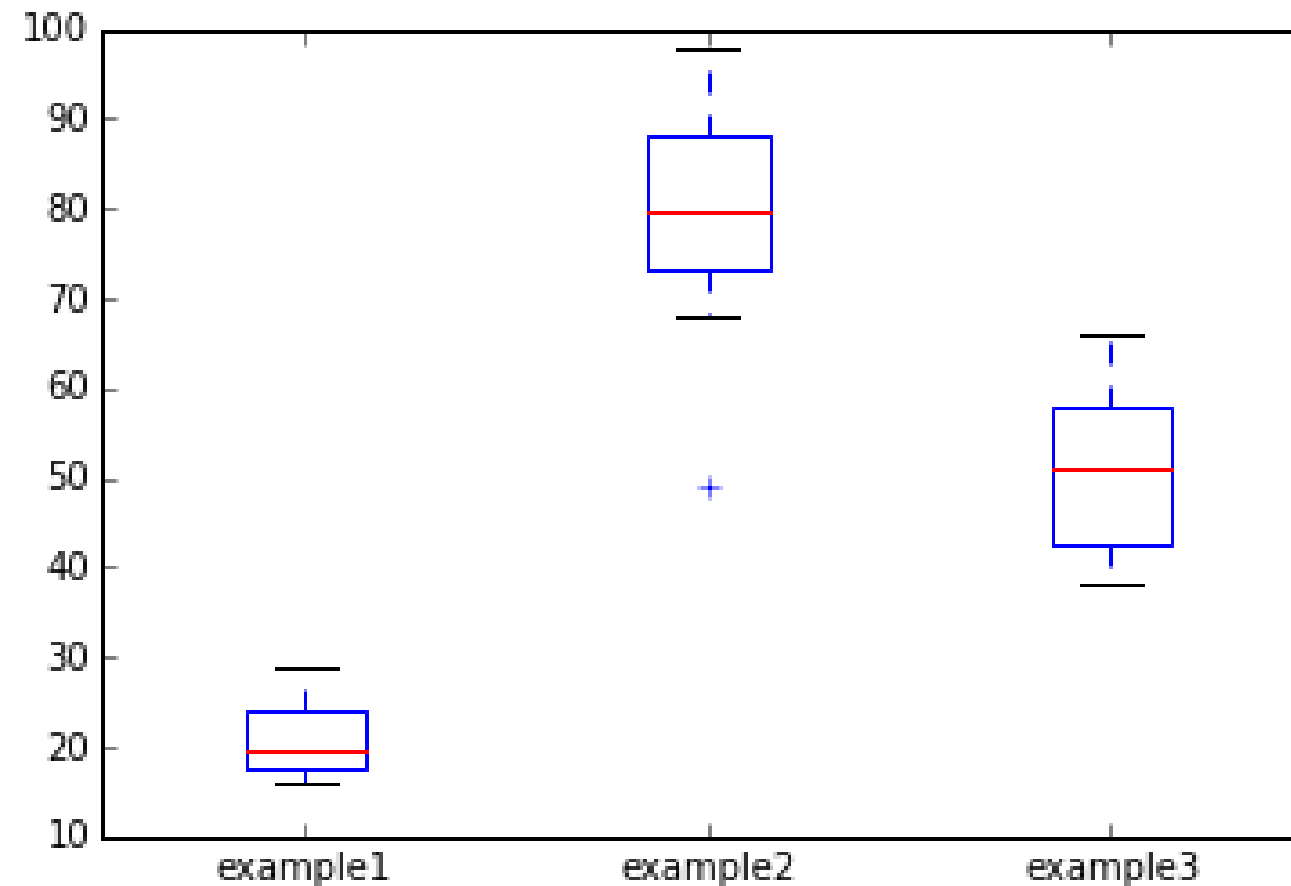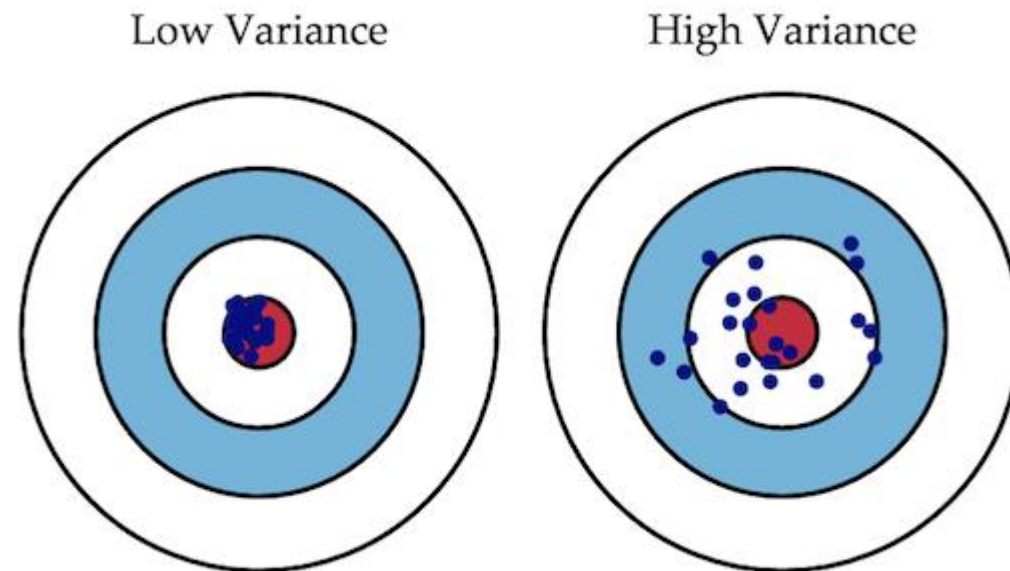
# BIAS VS. VARIANCE

‣ Error due to **bias** is calculated at the difference between the *expected prediction* of our model and the *correct value* we are trying to predict.

‣ Imagine creating multiple models on various datasets. **Bias** measures *how far off in general* models' predictions are from the correct value.
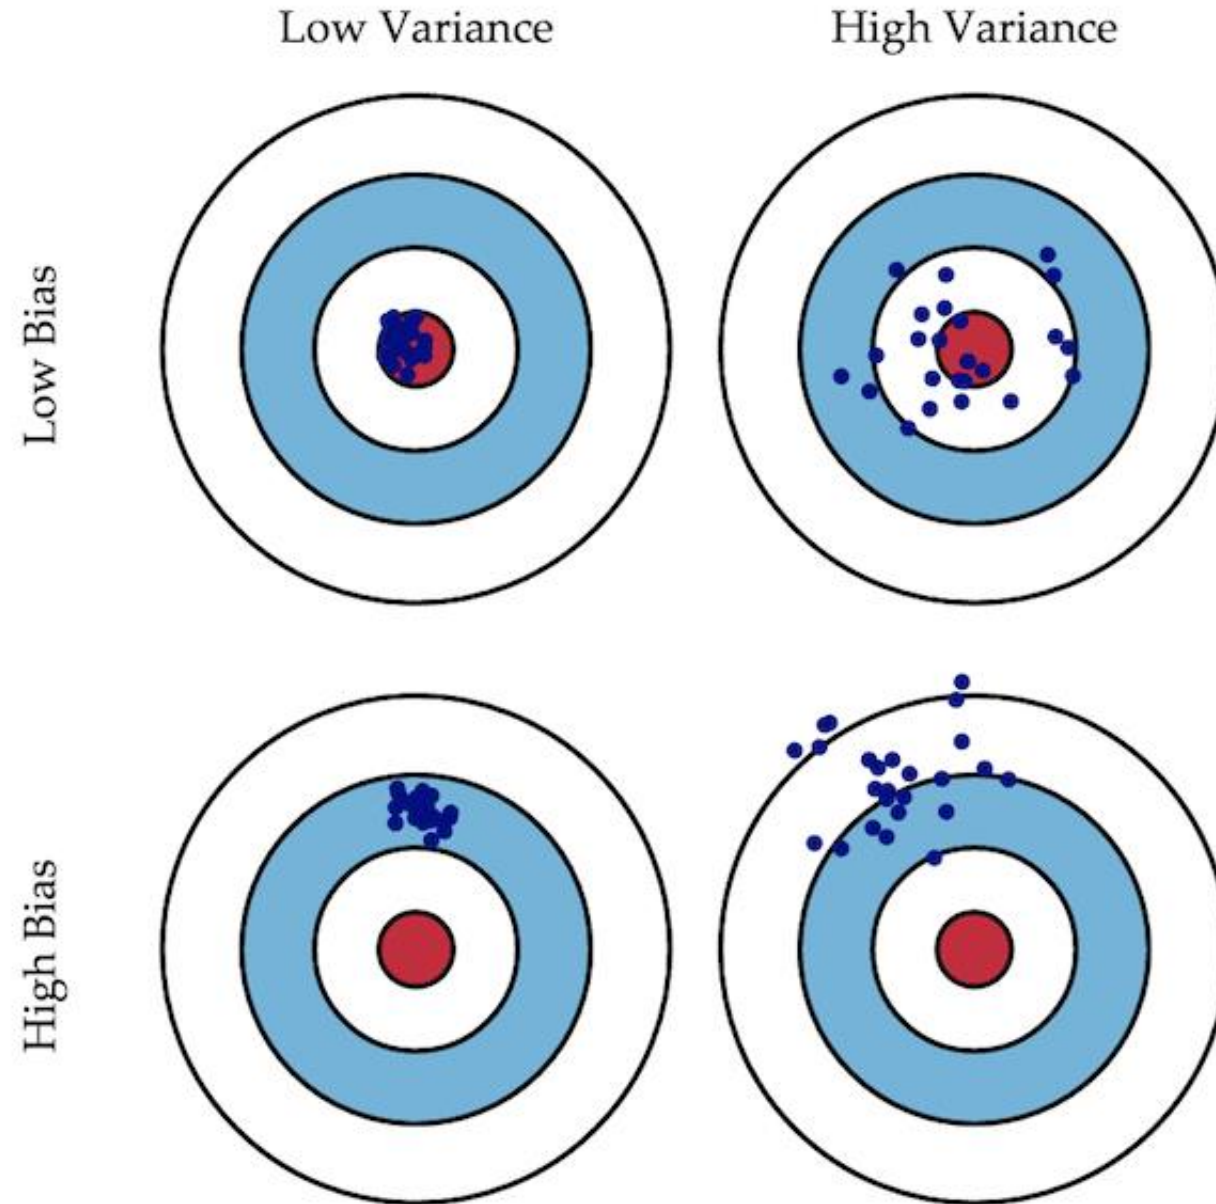


Low Bias

High Bias
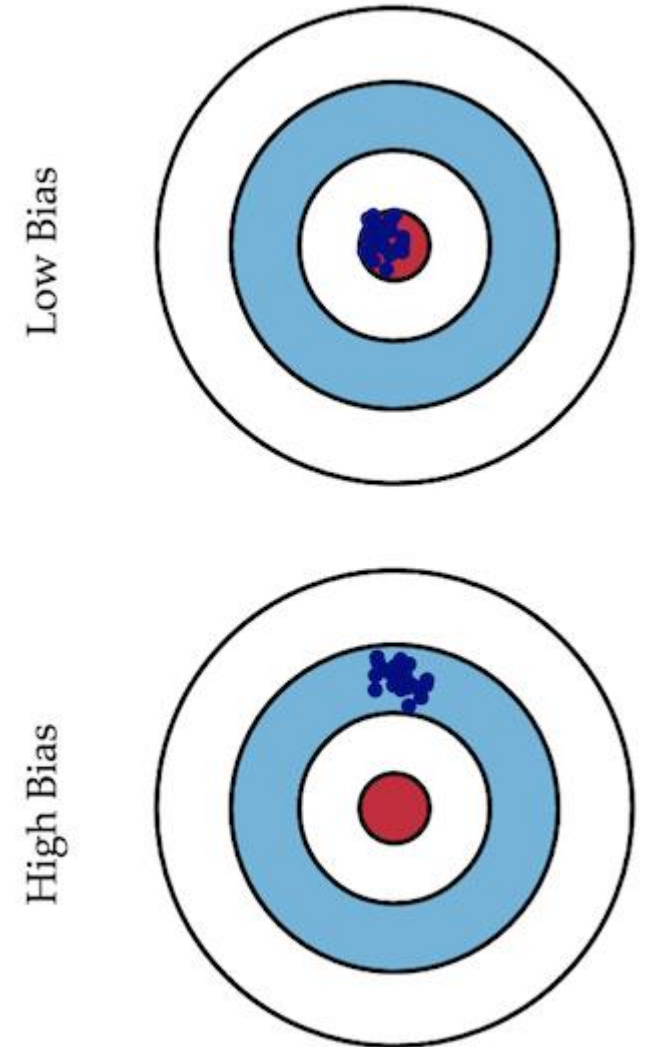
# BIAS VS. VARIANCE

‣ Error due to **variance** is taken as the variability of a model prediction for a given point.

‣ Imagine creating multiple models on various datasets. The **variance** is *how much the predictions for a given point vary* between different realizations of the model.


Low Variance          High Variance

# BIAS VS. VARIANCE

# STANDARD DEVIATION

‣ Standard deviation (SD, σ for population, s for sample) is a measure that is used to quantify the amount of variation or dispersion of a set of data values.

‣ Standard deviation is the square root of variance.

$$s = \sqrt{\frac{\sum(x - \bar{x})^2}{n - 1}}$$

# STANDARD ERROR

‣ The standard error of the mean (SEM) quantifies the precision of the mean.

‣ It is a measure of how far your sample mean is likely to be from the true population mean.

‣ It generally increases with the size of an estimate, meaning a large standard error may not indicate the estimate of the mean is unreliable.

‣ It's often better to compare the error in relation to the size of the estimate.

# STANDARD ERROR

$$SE_{\bar{x}} = \frac{s}{\sqrt{n}}$$

# CODEALONG PART 3: STANDARD DEVIATION & VARIANCE

‣ You can calculate variance and standard deviation easily in Pandas.
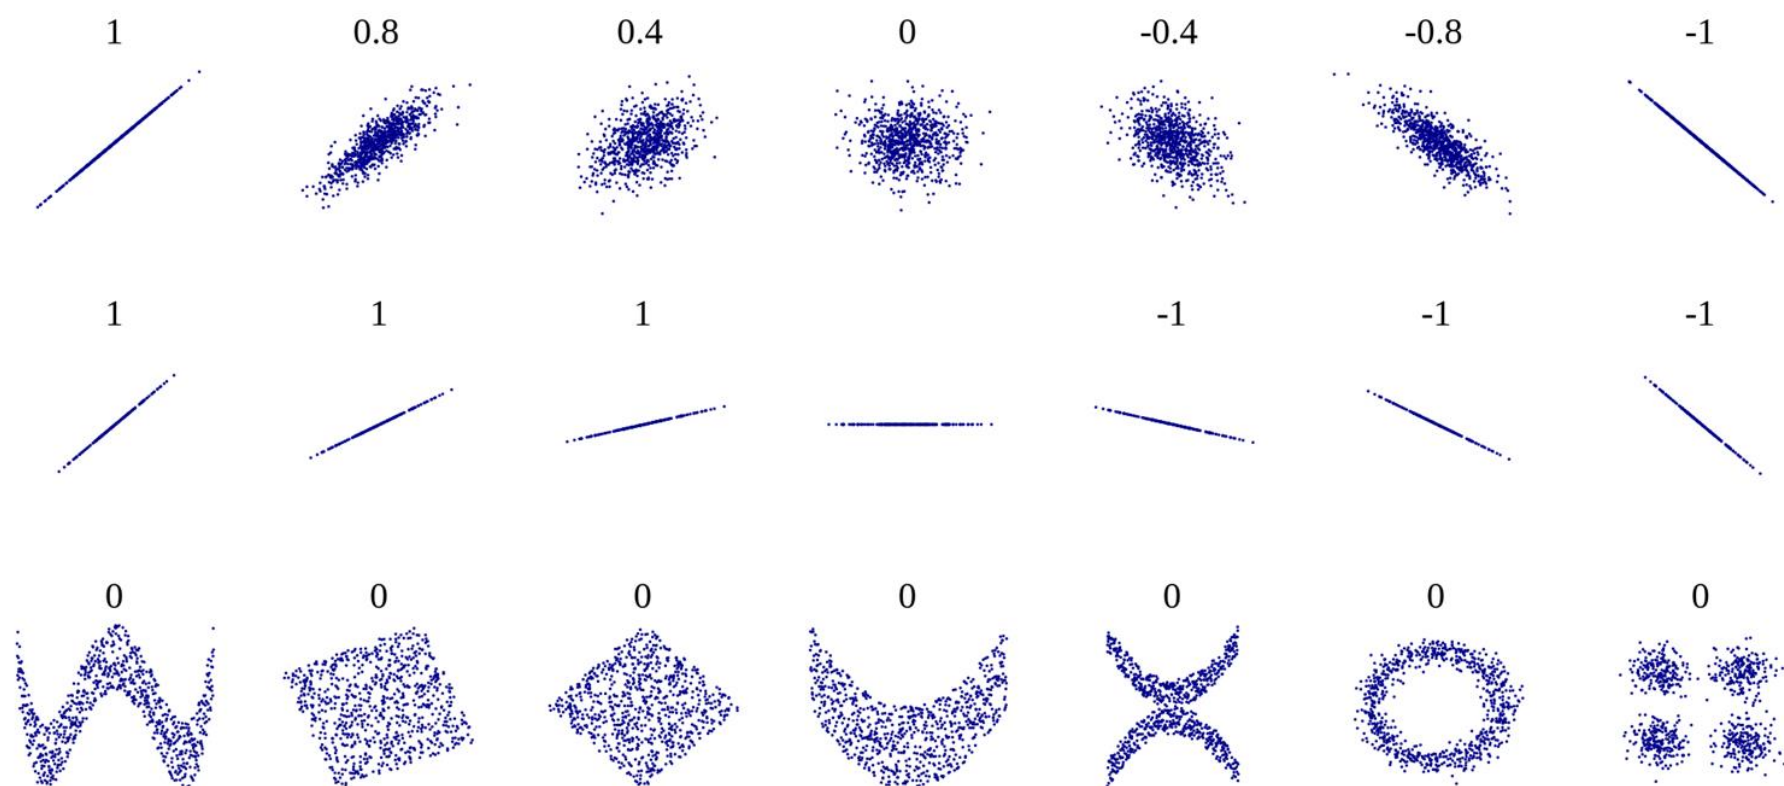
Methods include:

.std() - Compute Standard Deviation

.var() - Compute variance

.describe() - short cut that prints out count, mean, std, min,

quartiles, max

# CORRELATION

‣ The correlation measures the extent of interdependence of variable quantities.

‣ Example correlation values

# CONTEXT

‣ For most projects, descriptive stats will come first. These help you get to know your dataset better.

‣ Sometimes, descriptive stats may be all you need to answer your question.

# IS THIS NORMAL?

# THE NORMAL DISTRIBUTION

‣ A normal distribution is often a key assumption to many models.

‣ The normal distribution depends upon the *mean* and the *standard deviation*.

‣ The *mean* determines the center of the distribution. The *standard deviation* determines the height and width of the distribution.
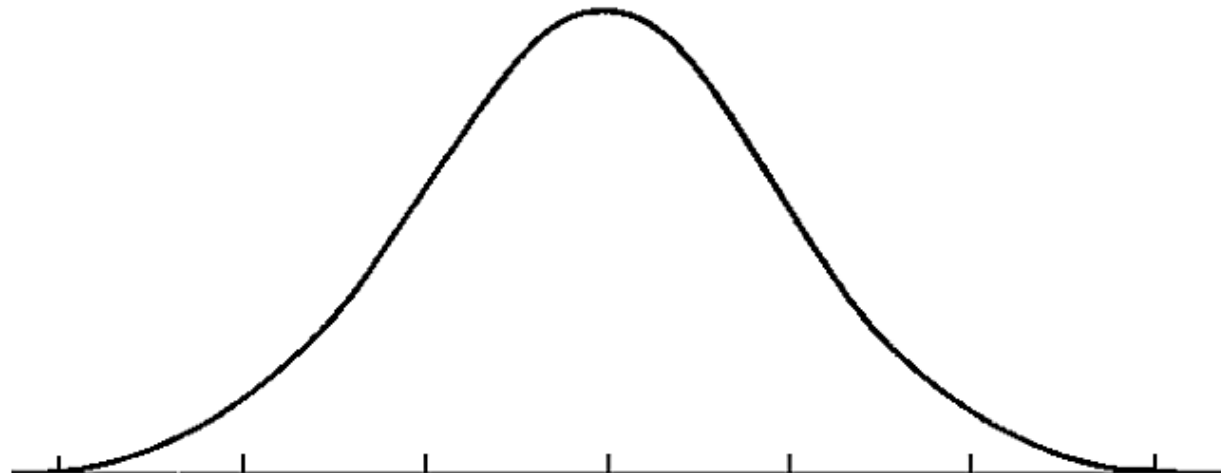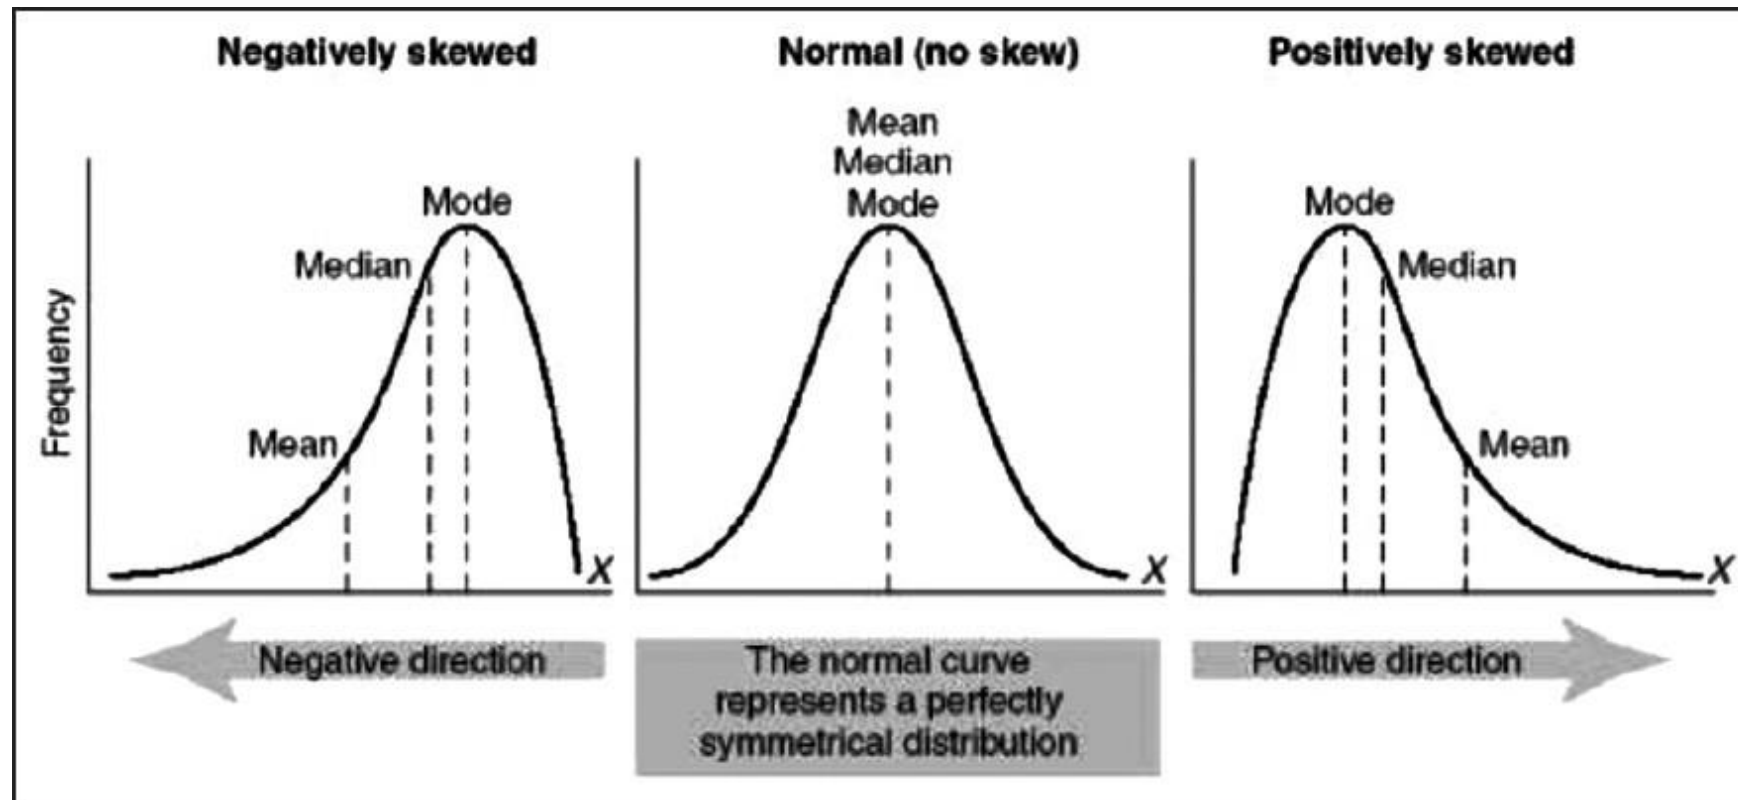
# THE NORMAL DISTRIBUTION

‣ Normal distributions are symmetric, bell-shaped curves.

‣ When the standard deviation is large, the curve is short and wide.

‣ When the standard deviation is small, the curve it tall and narrow.
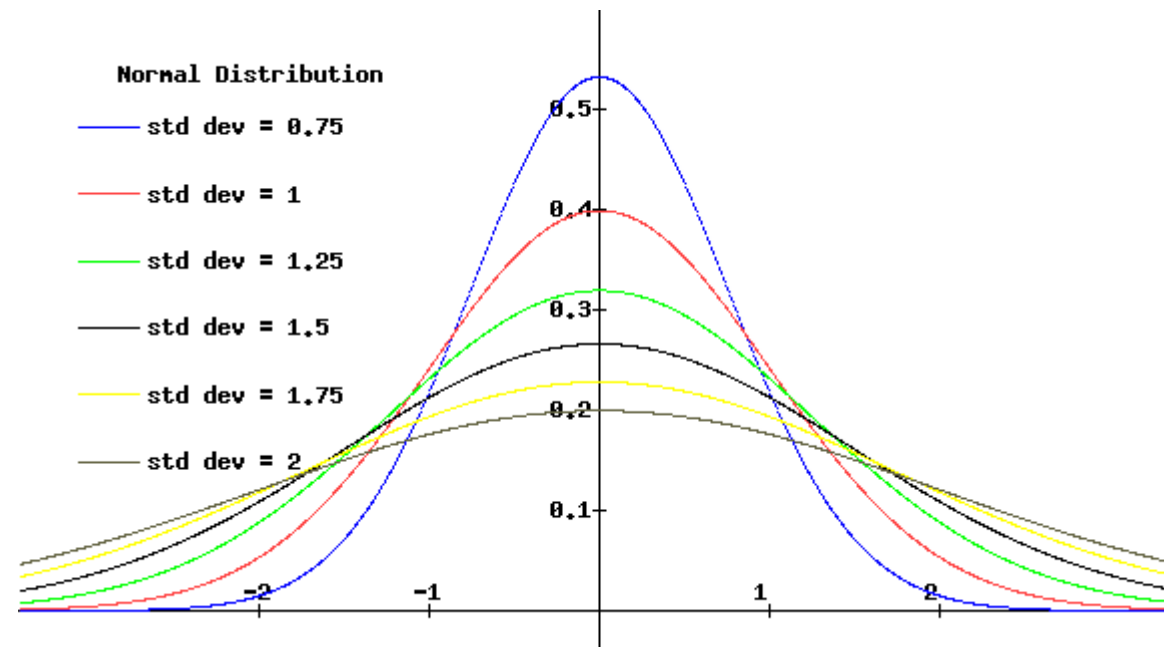
# SKEWNESS

▸ Skewness is a measure of the asymmetry of the distribution of a random variable about its mean.

▸ Skewness can be positive or negative, or even undefined.

# KURTOSIS

‣ Kurtosis is a measure of whether the data are peaked or flat relative to a normal distribution.

‣ Datasets with high kurtosis tend to have a distinct peak near the mean, decline rather rapidly, and have heavy tails.
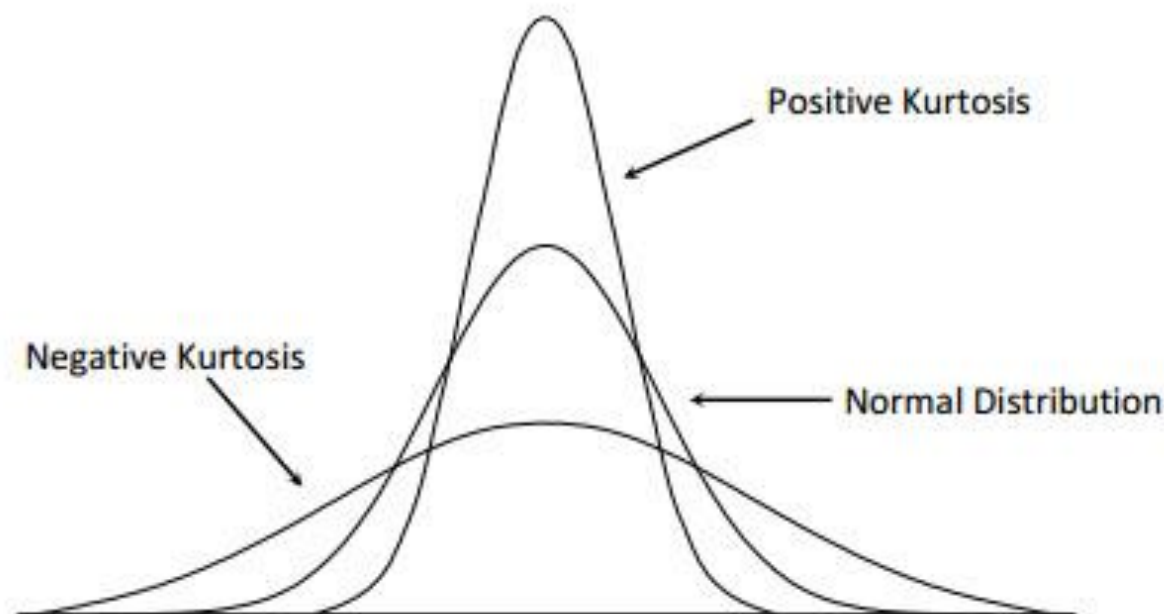
# DETERMINING THE DISTRIBUTION OF YOUR DATA

# DETERMINING THE DISTRIBUTION OF YOUR DATA

‣ Follow along as we walk through this in an iPython Notebook.

# VARIABLE TYPES

# VARIABLE TYPES

‣ Numeric variables can take on a large range of non-predetermined, quantitative values. These are things such as height, income, etc.

‣ Categorical variables can take on a specific set of variables. These are things such as race, gender, paint colors, movie titles, etc.

# CLASSES

# CLASS/DUMMY VARIABLES

‣ Let's say we have the categorical variable `area`, which takes on one of the following values: `rural`, `suburban`, and `urban`.

‣ We need to represent these numerically for a model. So how do we code them?

# CLASS/DUMMY VARIABLES

‣ How about `0=rural`, `1=suburban`, and `2=urban`?

# CLASS/DUMMY VARIABLES

‣ But this implies an ordered relationship - is `urban` twice `suburban`? That doesn't make sense.

‣ However, we can represent this information by converting the one `area` variable into two new variables:

`area_urban` and `area_suburban`.

# CLASS/DUMMY VARIABLES

‣ We'll draw out how categorical variables can be represented without implying order.

‣ First, let's choose a reference category. This will be our "base" category.

‣ It's often good to choose the category with the largest sample size and a criteria that will help model interpretation. If we are testing for a disease, the reference category would be people without the disease.

# CLASS/DUMMY VARIABLES

‣ Step 1:  Select a reference category.  We'll choose `rural` as our reference category.

‣ Step 2:  Convert the values `urban`, `suburban`, and `urban` into a numeric representation that does not imply order.

‣ Step 3:  Create two new variables:  `area_urban` and `area_suburban`.

# CLASS/DUMMY VARIABLES

‣ Why do we need only two dummy variables?

| rural | urban | suburban |
|-------|-------|----------|

‣ We can derive all of the possible values from these two. If an `area` isn't `urban` or `suburban`, we know it must be `rural`.

‣ In general, if you have a categorical feature with k categories, you need to create k-1 dummy variable to represent all of the information.

# CLASS/DUMMY VARIABLES

‣ Let's see our dummy variables.

|  | **area_urban** | **area_suburban** |
|---|---|---|
| rural | 0 | 0 |
| suburban | 0 | 1 |
| urban | 1 | 0 |

‣ As mentioned before, if we know `area_urban=0` and `area_suburban=0`, then the area must be `rural`.

# CLASS/DUMMY VARIABLES

‣ We can do this for a `gender` variable with two categories: `male` and `female`.

‣ How many dummy variables need to be created?

# CLASS/DUMMY VARIABLES

‣ # of categories - 1 = 2 -1 = 1

# CLASS/DUMMY VARIABLES

‣ We will make `female` our reference category. Thus, `female=0` and `male=1`.

|  | gender_male |
|---|---|
| female | 0 |
| male | 1 |

‣ This can be done in Pandas with the `get_dummies` method.

# DUMMY COLORS

# ACTIVITY: DUMMY COLORS

DIRECTIONS (15 minutes)

It's important to understand the concept before we use the Pandas function `get_dummies` to create dummy variables.  So today, we'll create our dummy variables by hand.

1.  Draw a table like the one on the white board.
2.  Create dummy variables for the variable "colors" that has 6 categories:  blue, red, green, purple, grey, and brown.  Use grey as the reference.

EXERCISE

DELIVERABLE

Dummy variables table for colors

# TOPIC REVIEW

# REVIEW

‣ Let's go through the process for creating dummy variables for "colors".

   ‣ We talked about several different types of summary statistics, what are they?

   ‣ We covered several different types of visualizations; which ones?

   ‣ We talked about the normal distribution; how do we determine your data's distribution?

‣ Any other questions?

## ASSIGNMENT 1: Data Exploration

# DUE: 3/27 @ Midnight

Project: Using a dataset of your choice, create a Jupyter Notebook or python script that demonstrates your understanding of python, pandas, and statistics (this weeks topic).

- ‣ Requirements:
  - ‣ Import and clean data
  - ‣ Create calculated columns (e.g. dummy variables, new groupings)
  - ‣ Describe your data (statistics, visualizations)
  - ‣ Identify areas for further research
  - ‣ Bonus: Combine with a second dataset

# EXIT TICKET

TICKET LINK: http://bit.ly/2nxjvB2

Lesson: 3
Topic: Statistics Fundamentals - 1